

IGARSS 2026

Earth Embeddings as Products

Taxonomy, Ecosystem, and Standardized Access

Heng Fang* — KTH Royal Institute of Technology

Adam J. Stewart* — Technical University of Munich

Isaac Corley — Taylor Geospatial


Xiao Xiang Zhu — Technical University of Munich

Hossein Azizpour — KTH Royal Institute of Technology

*equal contribution · since expanded into the book chapter *Earth Embeddings* (arXiv, August 2026)

Run the model once, reuse the vectors

NASA's EOSDIS alone holds **178.7 PB** of imagery and grows by 160 TB per day. Every team downloads the same pixels again and repeats the same preprocessing and GPU inference. Embedding products run the model once and distribute the vectors as reusable data.



The Earth is one large document

Embedding retrieval matured in image search and in document retrieval for LLMs (RAG). The recipe is to chunk, embed, index, and retrieve. The Earth archive is one enormous document.


But there is no obvious chunk:

What is the **spatial extent** of a chunk?


One **timestamp**, or a time series?

What is the **resolution** of each point in it?

Which **signal** (MSI band, SAR, elevation)?



Three families of Earth embeddings



Patch products summarize km-scale tiles

All known patch embedding products as of July 2026. *sparse spatial or temporal coverage.

Product	Extent	Resolution	Years	Dim.	Dtype
MOSAIKS (2021–25)	USA / Global	1 km / 0.01°	2018 / 2019	8192 / 4000	float32
Clay v0 / v1.5 (2024–26)	USA / Global	154 m – 5.12 km	2010–2025*	768–1024	float32
Major TOM (2024–26)	Global*	120 m – 3.84 km	2016–2024*	256–2048	float32
Earth Index (2025)	Global	320 m	2024	384	float32
Copernicus-Embed (2025)	Global	0.25°	2021	768	float32

Major TOM is the largest family, with more than a dozen model variants on one shared grid. Every product stores float32. One vector summarizes an entire mosaic tile, which suits retrieval more than dense mapping.

Pixel products store a vector per pixel-year

All known pixel embedding products as of July 2026. Every product is built from annual time series. *sparse coverage.

Product	Extent	Resolution	Years	Dim.	Dtype
Presto Embeddings (2025)	Togo	10 m	2019–2020	128	uint16
Tessera Embeddings (2025)	Global*	10 m	2017–2025*	128	int8 → float32
Google Satellite Embedding (2025)	Global	10 m	2017–2025	64	int8 → float64
Embedded Seamless Data (2026)	Global	30 m	2000–2024	12	uint16 → float32

Google Satellite Embedding covers the full Sentinel era at 64 dimensions. **Embedded Seamless Data** trades resolution for 25 years of Landsat history. Storage pressure drives the int8 and uint16 dtypes.

Industry is selling embeddings as the product

Descartes Labs 2017

GeoVisual Search was the first planet-scale visual search over satellite imagery features. A query patch of a wind turbine returns wind turbines worldwide. It predates the current products by eight years.

Earth Genome 2025

Earth Index packages Clay embeddings on Source Cooperative into a search service for investigative journalists and conservation teams, who use it to find mining sites, airstrips, and feedlots.

LGND 2025

Raised **\$9M** in 2025 to build services on geographic embeddings, and published the full global Clay v1.5 float32 corpus on Source Cooperative.


All three build their products on the embeddings rather than the model weights.

Formats, grids, and hosting differ per product

- Products are scattered across **Source Cooperative** (Clay, Earth Index), Hugging Face (Major TOM), Earth Engine (Google, Presto), and private servers (Tessera).
- Formats span **GeoParquet**, GeoTIFF with implicit CRS assumptions, and raw NumPy arrays with no CRS or bounds ("numbers and a prayer").
- Each product defines its own **tiling grid**, so any cross-product comparison starts with reprojection.
- **One flipped coordinate** in the Google rasters forced fixes in **GDAL, rasterio, and TorchGeo**.

Every team solves distribution on its own, so every user pays an integration cost for every product.

Storage costs span five orders of magnitude



Patch products stay in the MB–GB range, while dense 10 m pixel products reach tens of TB. A single full download of Presto’s Africa coverage adds ~**\$6.9k in egress** on top of storage.

Hosting and format choices decide usability

- **Hugging Face** enforces storage caps and API rate limits, so bulk pulls of TB-scale products throttle or fail.
- **Tessera** distributes `.npy` tiles plus a metadata sidecar. NumPy has no HTTP range requests, so reads pull whole tiles that a **COG or Zarr would stream**. It is GeoTIFF without the streaming or the metadata.
- The **AlphaEarth embeddings** were available only through Earth Engine or a requester-pays bucket, where the reader pays the egress. **Taylor Geospatial rehosted them in a non-requester-pays bucket on Source Cooperative**, with free egress, HTTP range reads, and CORS for browser streaming.
- Several products are so **sparse in space and time** that no common footprint exists for comparison.

No product is fully reproducible

- Clay, Earth Index, and Copernicus-Embed release code, weights, and data under permissive licenses.
- Tessera releases code, weights, and embeddings openly, but records **no metadata about which inputs built each tile**, so its outputs cannot be audited.
- Major TOM's CC-BY-SA pretraining data makes its embeddings copyleft, deterring commercial users.
- AlphaEarth and ESDNet keep code and weights proprietary. The embeddings are CC-BY, but no one outside can regenerate them.
- No product provides checksums. Archives keep reprocessing scenes, so exact inputs are unrecoverable.

A foundation model is not an embedding product

Foundation models (dynamic inference)

Distributed as public weights (DOFA, OlmoEarth). Users manage preprocessing and GPU inference themselves. This is flexible but raises a hardware and engineering barrier.

Embedding products (static data)

Distributed as pre-computed vector archives (Google Satellite Embedding). The assets are frozen and versioned, so analysis proceeds without the model or its compute.

A product is pinned to the data snapshot it was computed on, so **results measured on a product do not generalize to the model behind it.**

Paper metrics do not reproduce from the products

The AlphaEarth and Tessera papers report benchmark numbers from internal pipelines. The same tasks, evaluated on the **released annual embedding products**, give different numbers.

- The papers embed **exact input stacks**. The products are **annual composites** over reprocessed archives.
- Benchmark numbers should describe the **downloadable data** rather than the model.

The model is not the map

The Model is Not the Map

Please Look at the Data



CHRISTOPHER REN

JUN 21, 2025

5

2

1


Share

TLDR;


- Make maps, not war (just) benchmarks
- I made the Santa Fe Quartet as a geospatial equivalent to Anscombe's quartet
- If you are interested in building this kind of map-benchmark let's chat!

Ren corrupts one land cover map four ways. Every version scores **F1 $\approx$ 0.73–0.74**, but each error is obvious on the map. Release predictions and embeddings alongside checkpoints. Nobody knows a model better than the team that trained it.


a) Random Salt




b) Neighbor Dilation



c) Tiling Issue



d) Class Confusion



Embeddings improve mapping, not spatial transfer

Task	Embeddings	Finding
Cropland mapping, Togo	Presto, Google	Presto + Random Forest gives the best F1
Tree species, Dutch forest inventory	Presto, Google, Tessera	+2–9 points over hand-designed time-series features
Landslide susceptibility, TW·HK·IT	Google	64-d embeddings beat conventional conditioning factors
Poverty mapping, Sub-Saharan Africa	Google + GNN	large storage/preprocessing savings vs. raw Sentinel-2
Scene classification, EuroSAT-Embed	Google, Tessera, OlmoEarth	pooling choice cuts the geographic gap by >50%
Fusion across six tasks	Google, Tessera, GeoCLIP, SatCLIP	fused embeddings beat the best single model in 4 of 6

Performance drops under **spatial transfer**, and annual composites wash out sub-annual dynamics. Validate with geographic splits and simple baselines before trusting any product.

No one knows which product is best for your task

An audit of 152 geospatial foundation model papers found **46 cross-paper disagreements of ≥ 10 points** for the same model, benchmark, and protocol.

94 of 126 papers use a pretraining configuration that appears in no other paper. **39% release no weights**, so their results cannot be re-run at all.

We built **torchgeo-bench**, a maintained harness for frozen backbones with shared datasets, consistent probes, and bootstrapped confidence intervals.



Products are complementary

You do not have to pick one product.

Fusing AlphaEarth, Tessera, GeoCLIP, and SatCLIP beats the best single product on **4 of 6 downstream tasks**.

Which products complement each other depends on the task and region, so a probe trained on the concatenated embeddings is a cheap first experiment.

BETTER TOGETHER: EVALUATING THE COMPLEMENTARITY OF EARTH EMBEDDING MODELS

A PREPRINT

Thijs I. van der Plas¹ **Jacob JW Bakermans²** **Vishal Nedungadi¹**
Gabriele Tjiauitjia¹ **Marc Ruthwurm¹** **Ioannis N Athanasiadis¹**
¹Wageningen University, NL ²University College London, UK
 {thijs.vanderplas, ioannis.athanasiadis}@wur.nl

May 19, 2026

ABSTRACT

Earth embedding models transform Earth observation data into embeddings uniquely tied to locations on the Earth's surface. These models are typically evaluated in isolation, comparing the downstream task performance across different Earth embeddings. However, spatially aligned embeddings can naturally be fused, providing richer information per location, a capability that isolated evaluations fail to capture. We therefore propose assessing Earth embeddings by their *complementarity*: the performance gain of fused embeddings over the best single-model baseline. To operationalise this, we introduce an embedding complementarity index applicable to any embedding and task, and evaluate four Earth embedding models (AlphaEarth, Tessera, GeoCLIP, SatCLIP) in isolation, in all pairs, and jointly across six downstream tasks. Fused embeddings outperform the best single model in four out of six tasks, confirming that single-embedding evaluations often underestimate Earth embedding capabilities. Complementarity proves both task- and location-dependent. Further, for a land cover regression task, we find that complementarity is partially determined by the spatial scale of land cover classes. Complementarity reframes Earth embeddings: the greatest future gains may come not from any single Earth embedding model, but from combinations that are better together.


1 Introduction

Earth embedding models construct compressed embeddings from rich, multimodal Earth observation data, which can be used for a variety of geospatial downstream tasks (Klemmer et al., 2025b). Compared to foundation model embeddings from other domains, such as vision or language, Earth embeddings have the unique property that any embedding is inherently linked to a location on the Earth's surface. This trait has turned Earth embeddings into commodities, which are pre-computed once and then retrieved by coordinates alone, bypassing the need for model inference (Fang et al., 2026). As a result of this ease of access, pre-computed Earth embeddings have now been used for a range of downstream tasks in agriculture, environmental sciences and ecology (Zvonkov et al., 2025; Ma et al., 2026; Young et al., 2026; Lisiasus et al., 2026; Ball et al., 2026). But while the *differences* in task performance between Earth embeddings are regularly assessed (Zvonkov et al., 2025; Lisiasus et al., 2026; Ball et al., 2026), the *complementarity* between different Earth embeddings is not.

This lack of complementarity assessments might stem from the notion that Earth embedding models are expected to encode all possible information about a location within a single vector. Yet, Earth embedding models are developed with varying training objectives, input data modalities and training datasets. For example, explicit geospatial foundation models map Earth observation images to pixels embeddings, and are trained using masked auto-encoders (Brown et al., 2025) or pixel-wise contrastive loss (Tseng et al., 2023; Feng et al., 2025), while implicit neural representation models map coordinates to location embeddings via various contrastive loss objectives (Vivanco et al., 2023; Klemmer et al., 2025a; Dhakal et al., 2025). This variety of Earth embeddings suggests that each might capture a different "view" of the same location that is partially unique, and consequently they can be complementary sources of information. To test this,


Smaller, sliceable vectors are underexplored

- **Quantization.** Google and Tesseract use int8 at no measurable cost. No product goes lower, yet binary recovers ~65% of float32 nearest neighbors.
- **Disentangled representations.** VAE-style training gives each dimension a separate meaning. Untried for Earth embeddings.
- **Matryoshka learning.** Tesseract v2, Clay, and our MIND location encoder train nested dimensions, so users truncate to their budget. No other products do.



the leading dimensions form a usable embedding at any width

Compressed embeddings run in the browser



TerraBit — 50M Clay v1.5 patches binarized to 128 bytes each and streamed from Source Cooperative. Hamming search runs in a Web Worker with no backend.

isaac.earth/terrabit



DeltaBit — AlphaEarth pixel differences at 8 bytes per pixel (PCA-8 + int8) as XYZ GeoTIFF tiles. The user labels, trains, and maps change in the browser (Seattle, 2020 → 2024).

calebrob.com/deltabit

Streamable formats, embedded metadata, and int8

Embedding type	Format
Location / implicit	ONNX or PyTorch
Patch / region	GeoParquet
Pixel, snapshot	GeoTIFF or GeoZarr
Pixel, time series	GeoZarr

- Provide a **model card**: sensors, temporal window, CRS, grid, dtype, quantization, license.
- Store metadata **inside the files** rather than in separate docs.
- **int8 by default**; PCA to 64 dims costs <2% accuracy for 64× compression.
- **Runnable benchmarks** instead of leaderboards.

Takeaways and open problems

Pick the family by task: implicit for location context, patch for retrieval, pixel for dense mapping. Validate under geographic splits against simple baselines, and consider fusing products instead of picking one. Open problems are standardized formats and provenance, ocean and atmosphere coverage, uncertainty layers, and shared benchmarks (identical models differ by more than ten points across papers).

Chapter arxiv.org/abs/2608.03410

Survey github.com/hfangcat/Awesome-Geospatial-Embeddings

Bench torchgeo.org/torchgeo-bench

Contact isaac.corley@taylorgeospatial.org



these slides